

Aarya Vaidya

AI/ML Engineer (Final-year B.Tech, AI/ML 2027) | Production LLM Applications · RAG · Agentic AI Systems

+91 8149904590 · aaryasharadvaidya.10@gmail.com · linkedin.com/in/aarya-vaidya-013227211 · github.com/rajaaryan779 · Pune, India
aarya-portfolio-kohl.vercel.app

Available for remote roles now · open to relocation and visa sponsorship · full-time from mid-2027

Summary

Final-year **B.Tech AI/ML** student (graduating 2027) who ships production LLM systems end to end. At **OyeLabs Technologies** I built a deterministic chat fast-path for **Jenna** (AI financial coaching) that **cut p50 latency ~17x**, plus a **fabrication guard** that grounds every dollar figure against real transaction data before it is sent, and **Market Expansion Engine**, a multi-tenant B2B outreach SaaS secured by Postgres **Row-Level Security**. Independently shipped **Ai-Prashashan**, a retrieval-grounded Nemotron assistant (Kotlin/Android) that answers only from official source text and refuses when unsupported. A recurring theme in my work: making LLMs **refuse rather than fabricate**. Comfortable across the stack — **FastAPI**, **RAG**, multi-agent orchestration with **LangGraph**, vector search (pgvector, FAISS, Pinecone), and **AWS** — with daily use of agentic dev tooling (Cursor, Claude Code). Seeking a graduate or junior AI/ML engineering role.

Experience & Internships

OyeLabs Technologies

Remote, India

AI Engineering Intern

Jun 2026 to Present

- Building production **AI automation** systems across backend, frontend, database, and mobile app development; shipped AI features now running in **Jenna** and **WhatsCRM** (see Projects).
- Built **Market Expansion Engine (MEE)**, a **multi-tenant** B2B lead-generation and AI cold-outreach platform (Next.js · Supabase RLS · n8n · Groq) that scrapes, enriches, scores, and reaches out to leads end to end with tenant isolation enforced at the database layer.

Elevate Labs

Remote, India

AI Intern

2025

- Built an **LSTM** stock-price forecasting pipeline (TensorFlow/Keras) with engineered features (moving average, RSI, lag) and a backtesting harness, outperforming a linear baseline.

Projects

Jenna - AI Financial Coaching App

FastAPI · Amazon Bedrock (Claude) · PostgreSQL · Plaid · Auth0

OyeLabs · AWS dev

- Built a **deterministic chat fast-path** that classifies read-only summary intents and template-renders pre-computed figures with no LLM call, cutting p50 latency from ~1,950ms to ~115ms (**~17x**) by bypassing the multi-tool agent loop.
- Engineered a **fabrication guard** that verifies every dollar amount in an LLM reply against grounded transaction data and falls back when ungrounded, **eliminating hallucinated balances**; shipped to AWS dev via GitHub Actions (ECS, Alembic).

WhatsCRM - WhatsApp CRM

OpenAI Embeddings · Supabase pgvector · FastAPI · Webhooks

OyeLabs

- Built the ingestion pipeline (webhook to **OpenAI** text-embedding-3-small to Supabase **pgvector** to AI reply suggestions) with idempotent upserts and **zero message loss** under concurrent load, plus semantic search over prior customer conversations.

Ai-Prashashan - Grounded GR Assistant

Kotlin · Jetpack Compose · NVIDIA Nemotron 3 Ultra · Firebase · Cloudflare R2

github.com/rajaaryan779/Ai-Prashashan

- Marathi-language Android app for Maharashtra government employees: a **retrieval-grounded** Nemotron assistant answering service-rule questions over **440+ official Government Resolutions** across 94 topics, ranked by **TF-IDF/IDF** with a Marathi-aware tokenizer (Devanagari-digit and ordinal normalization).
- Engineered a **strict grounding guard** (refuse-when-unsupported prompt contract, temperature 0.15) so every reply is drawn only from retrieved GR text and cites its source PDF, **eliminating hallucinated service rules** when retrieval finds no support.
- Built the content pipeline: reverse-engineered the state GR portal's ASP.NET WebForms flow (connection-pinned VIEWSTATE) to bulk-download **534 GRs**, then Tesseract OCR (Marathi + English) to extract searchable text into Cloudflare R2 for retrieval.

Market Expansion Engine (MEE)

Next.js 14 · TypeScript · Supabase (Postgres + RLS) · n8n · Groq (Llama 3.3 70B)

github.com/rajaaryan779/MEE · B2B Lead-Gen & AI Outreach

- Converted a single-tenant app into a **multi-tenant SaaS**: composite (owner_id, place_id) keys and **Row-Level Security across all 8 tables** with a security-invoker view, so tenant isolation is enforced by Postgres — not hopeful application code.
- Designed a thin Next.js control plane that authenticates and stamps owner_id, offloading long-running scrape / crawl / LLM work to **n8n** past serverless timeouts, with a **live dashboard via Supabase Realtime** and lead-quality scoring as an IMMUTABLE Postgres function.

AutoStream Agent

LangGraph · Gemini · Python

github.com/rajaaryan779/autostream_agent

- Social-to-lead qualification agent (ServiceHive): a 4-node **LangGraph** state machine (detect intent, collect, extract lead, respond) with keyword-grounded retrieval over a local knowledge base, typed state, and structured lead extraction.

Binance Futures Trading Bot

Python · Binance API · REST

github.com/rajaaryan779/trading_bot

- Testnet order CLI (PrimetradeAI) with **HMAC-SHA256** signed REST calls, input validation, structured logging, and modular market/limit order execution.

AEO Diagnostic Tool

FastAPI · Gemini 2.5 Flash · OpenRouter · Vercel

github.com/rajaaryan779/aeo-diagnostic

- Concurrent multi-LLM engine querying 3 models simultaneously to score product AI-search visibility (0 to 100 AEO score). **Shortlisted Top 25 for the Pixii.ai founding engineer role.**

Grid07 - Cognitive Routing & RAG Engine

LangGraph · FAISS · Groq · Python

github.com/rajaaryan779

- 3-stage pipeline: FAISS-based persona router to a LangGraph content engine with **schema-validated JSON output** to a retrieval stage with **prompt-injection input filtering**.

DreamHome

FastAPI · Next.js · PostgreSQL · Leaflet · JWT

github.com/rajaaryan779/Real-Estate-software

- Full-stack real estate platform with interactive maps. **JWT token caching cut auth validation from 15ms to <1ms** (93% reduction). Alembic zero-downtime migrations.

AI Automation Suite (19 Workflows)

n8n · LangChain · Pinecone · Whisper · Gemini

github.com/rajaaryan779/n8n-workflows

- RAG pipeline (Google Drive to Pinecone to Gemini), multi-agent bots with **per-user memory namespacing**, WhatsApp/Telegram integrations, and Whisper voice transcription.

Rlook - AI Virtual Try-On (In Development)

Next.js 15 · Three.js · MediaPipe · FastAPI

github.com/rajaaryan779

- Real-time 3D body segmentation and clothes overlay in-browser using **WebGL and body-pose tracking**, delivered fully in-browser with no native app or plugin.

Skills

AI & LLM:

LangChain · LangGraph · RAG · Grounded Generation & Hallucination Guardrails · FAISS · Pinecone · pgvector · Sentence-Transformers · Prompt Engineering · Agentic Workflows

LLM Models:

Claude · **GPT** · **Gemini** · LLaMA · Mistral · Whisper | served via Amazon Bedrock, OpenAI API, Groq & OpenRouter

Languages & Frameworks:

Python · TypeScript · JavaScript · **Kotlin** · SQL · FastAPI · Flask · Next.js · React · **Jetpack Compose (Android)** · REST APIs

Data & Infra:

PostgreSQL · AWS RDS · SQLAlchemy · Alembic · Supabase · **Firestore** · **Cloudflare R2** · Auth0 · JWT · asyncio · n8n · Webhooks

DevOps & Tools:

AWS (ECS, CDK, CloudWatch) · Docker · GitHub Actions · Vercel · Git · Linux · Gradle · Tesseract OCR · Cursor · Claude Code

Education

Pimpri Chinchwad University

Pune, India

B.Tech, Computer Science Engineering (AI & ML)

2024 to expected May 2027

Final year · 2nd Prize, PCU InnovateX (Engineers' Day 2025)

K.K. Wagh Polytechnic

Nashik, India

Diploma, Artificial Intelligence and Machine Learning

2021 to 2024

Certifications

Develop Generative AI Applications: Get Started - IBM | Coursera (2026), course 1 of 10 - RAG & Agentic AI Professional Cert

coursera.org/verify/KTAVNVHE8H6H

Claude with Amazon Bedrock - Anthropic (2026) verify.skilljar.com/c/mi62wwe5n6hi

Red Hat OpenStack Administration I - Red Hat (2026) credly.com/badges/3115833a-4a07-4cf9-ac22-855f007a5914

Foundations of Cybersecurity - Google | Coursera (2025) coursera.org/verify/MGIHIZWCAA54

Django Web Framework - Meta | Coursera (2025) coursera.org/verify/2Z1HXMV496N

Additional (Udemy): Python & Generative AI · AI Automation with n8n · Deep Learning & CNN · Flask API Development